

## Markov Decision Processes: General Model


Mohammad Sadegh Talebi  
[m.shahi@di.ku.dk](mailto:m.shahi@di.ku.dk)

Department of Computer Science



## Sequential Decision Making

Many tasks in real life are **online sequential decision-making** tasks that fall in the framework of **reinforcement learning**:



- Selling or buying an asset
- Inventory management
- Portfolio optimization
- Robotics
- Playing computer games
- Routing in networks
- Precision Agriculture and Farming
- LLMs




## Sequential Decision Making: General Setting

Almost all RL systems try to solve underlying **decision process**.

Minimal ingredients of a decision process:

- A notion of **state** capturing different situations
- **Actions** capturing options available at any situation
- A **reward signal** indicating the quality of the action taken




**Goal:** To maximize an objective function, often defined in terms of rewards  
 $r_1, r_2, \dots$

- E.g., maximize  $\sum_{t=1}^N r_t$  or  $\sum_{t=1}^N \log(1 + r_t)$



## Example: Balancing Cart-pole



**Task:** Make the pole upright for as long as possible.

- Notions of state:
  - position  $x$
  - position and angle  $(x, \theta)$
  - position, angle, velocity, angular velocity  $(x, \dot{x}, \theta, \dot{\theta})$
- Action: Force  $F$
- Reward: 1 if  $\theta < \theta_{\text{th}}$ , else 0. (E.g.,  $\text{th} = 10 \text{ deg}$ )



## Sequential Decision Making: General Setting


We consider discrete time systems, where time is divided into slots of equal length.

At each time  $t = 1, 2, \dots, N$ , an agent interacts with an **unknown** environment

- observes state  $s_t$ ,
- chooses an action  $a_t$  from a given action set, using a control policy

$$a_t = \text{policy}(s_1, a_1, r_1, \dots, s_{t-1}, a_{t-1}, r_{t-1}),$$

- receives (random) reward  $r_t$ .



- **Goal:** To maximize a function of rewards  $r_1, r_2, \dots, r_N$
- Observations and rewards are generated by an **uncertain** and (potentially) **unknown** environment.



## Markov Property

Different decision processes differ mainly on how  $s_{t+1}$  and  $r_t$  are generated.

Under the **Markov property**,  $s_{t+1}$  and  $r_t$  only depend on  $s_t$  and  $a_t$ .

$$\begin{aligned}\mathbb{P}(s_{t+1} = s' \mid s_1, a_1, \dots, s_{t-1}, a_{t-1}, \textcolor{red}{s_t}, \textcolor{red}{a_t}) &= \mathbb{P}(s_{t+1} = s' \mid s_t, a_t) \\ \mathbb{P}(r_t \mid s_1, a_1, \dots, s_{t-1}, a_{t-1}, \textcolor{red}{s_t}, \textcolor{red}{a_t}) &= \mathbb{P}(r_t \mid s_t, a_t)\end{aligned}$$

Namely, conditioned on  $(s_t, a_t)$ , the process is independent from the past.

This property defines **Markov Decision Processes (MDPs)**.



# Markov Decision Processes



## Markov Decision Process

A finite **Markov Decision Process (MDP)** is a tuple  $M = (\mathcal{S}, \mathcal{A}, P, R)$ :

- **State-space  $\mathcal{S}$**  (with size  $S$ )
- **Action-space  $\mathcal{A}$**  (with size  $A$ )
- **Transition function  $P$** : Selecting  $a \in \mathcal{A}$  in  $s \in \mathcal{S}$  leads to a transition to  $s'$  with probability  $P(s'|s, a)$ .  $P(\cdot|s, a)$  is a probability distribution over  $\mathcal{S}$ , i.e.,

$$\sum_{s' \in \mathcal{S}} P(s'|s, a) = 1$$

- **Reward function  $R$** : Selecting  $a \in \mathcal{A}$  in  $s \in \mathcal{S}$  yields a reward  $r \sim R(s, a)$ .
- The action-space may generally be state-dependent; we use  $\mathcal{A}_s$  to denote the set of actions available in state  $s$ .
- In general,  $\mathcal{S}$  or  $\mathcal{A}$  could be finite, countably infinite, or continuous.




## Interaction with MDP

An **agent** interacts with the MDP for  $N$  rounds.

At each time step  $t$ :

- The agent observes the current state  $s_t$  and takes an action  $a_t \in \mathcal{A}$
- The environment (MDP) decides a reward  $r_t := r(s_t, a_t) \sim R(s_t, a_t)$  and a next state  $s_{t+1} \sim P(\cdot | s_t, a_t)$
- The agent receives  $r_t$  (any time in step  $t$  before start of  $t + 1$ )



This interaction produces a trajectory (or history)

$$h_t = (s_1, a_1, r_1, s_2, a_2, r_2, \dots, s_{t-1}, a_{t-1}, r_{t-1}, s_t)$$



# Markov Property

MDPs adhere to the **Markov property**.

- At each time  $t$ ,  $s_{t+1}$  and  $r_t$  only depend on  $s_t$  and  $a_t$ .
- More precisely,

$$\mathbb{P}(s_{t+1} = s' \mid s_1, a_1, \dots, s_{t-1}, a_{t-1}, \textcolor{red}{s_t}, \textcolor{red}{a_t}) = \underbrace{\mathbb{P}(s_{t+1} = s' \mid s_t, a_t)}_{=P(s'|s_t, a_t)}$$

$$R(s_1, a_1, \dots, s_{t-1}, a_{t-1}, \textcolor{red}{s_t}, \textcolor{red}{a_t}) = R(s_t, a_t)$$



## Classification of MDPs based on Horizon $N$

- **Finite-Horizon MDPs:**  $N < \infty$ , and the goal is to solve

$$\max_{\text{all strategies}} \mathbb{E} \left[ \sum_{t=1}^{N-1} r(s_t, a_t) + r(s_N) \right]$$

- **Infinite-Horizon Discounted MDPs:**  $N = \infty$ , and given discount factor  $\gamma \in (0, 1)$ , the goal is to solve

$$\max_{\text{all strategies}} \mathbb{E} \left[ \sum_{t=1}^{\infty} \gamma^{t-1} r(s_t, a_t) \right]$$

- **Infinite-Horizon Undiscounted MDPs (Average-Reward MDPs):**  $N = \infty$ , and the goal is to solve


$$\max_{\text{all strategies}} \lim_{N \rightarrow \infty} \frac{1}{N} \mathbb{E} \left[ \sum_{t=1}^N r(s_t, a_t) \right]$$



# MDP Examples



## Example: Routing



**Task:** Find the maximum-weight route between node 1) and destination (node 8).

Modeling as finite-horizon MDP:

- States: Nodes in the graph  $S = \{1, 2, \dots, 8\}$
- Actions: Outgoing edges at each state; e.g.,  $\mathcal{A}_2 = \{\text{go to 4}, \text{go to 5}\}$
- Deterministic transitions
- Rewards: Edge weights
- Time horizon  $N$ : any number greater than the maximum path length ( $N \geq 4$ )



## Example: Product Management

Suppose we receive an order for a given product with probability  $\alpha$ . We can either process all the unfilled orders or process no order.

- The cost per unfilled order per period is  $c > 0$ , and the setup cost to process unfilled order is  $K > 0$ .
- Assume that the total number of orders that can remain unfilled is  $n$ .

**Task:** Find an order processing strategy that has minimal expected cost.



## Example: Product Management

Modeling as a discounted MDP:

- **State Space:** Define the state as the number of unfilled orders at the beginning of each period  $\Rightarrow \mathcal{S} = \{0, 1, \dots, n\}$ .
- **Action Space:** For  $s \neq 0, n$ , we have  $\mathcal{A}_s = \{J, \bar{J}\}$ , where  $J$  = processing unfilled orders and  $\bar{J}$  = processing no order  $\Rightarrow \mathcal{A}_0 = \{\bar{J}\}$  and  $\mathcal{A}_n = \{J\}$ .
- **Reward Function:**


$$\begin{aligned} R(i, J) &= -K, & R(i, \bar{J}) &= -ci, & i &= 1, \dots, n-1, \\ R(0, \bar{J}) &= 0, & R(n, J) &= -K. \end{aligned}$$

- **Transition Function:**

$$\begin{aligned} P(0|i, J) &= 1 - \alpha, & P(1|i, J) &= \alpha, & i &= 1, 2, \dots, n-1, \\ P(i|i, \bar{J}) &= 1 - \alpha, & P(i+1|i, \bar{J}) &= \alpha, & i &= 1, 2, \dots, n-1, \\ P(0|n, J) &= 1 - \alpha, & P(1|n, J) &= \alpha, \\ P(0|0, \bar{J}) &= 1 - \alpha, & P(1|0, \bar{J}) &= \alpha. \end{aligned}$$



## Example: Grid-world





**Task:** Find the shortest path from to .

- A grid-world with  $S = 20$  states, and 4 actions ( $\rightarrow, \uparrow, \downarrow, \leftarrow$ ).
- E.g.,  $a = \uparrow$  yields: moving  $\uparrow$  (w.p. 0.7), no move (w.p. 0.1), or moving  $\rightarrow$  or  $\leftarrow$  (each w.p. 0.1)
- Reward is 1 in , else 0.
- Once in :
  - the agent may stay there forever (**one-shot task**), or
  - the agent may be teleported to (**continual task**)



## Example: Grid-world




**Task:** Find the shortest path from to .

- A grid-world with  $S = 20$  states, and 4 actions ( $\rightarrow, \uparrow, \downarrow, \leftarrow$ ).
- E.g.,  $a = \uparrow$  yields: moving  $\uparrow$  (w.p. 0.7), no move (w.p. 0.1), or moving  $\rightarrow$  or  $\leftarrow$  (each w.p. 0.1)
- Reward is 1 in , else 0.
- Once in :
  - the agent may stay there forever (**one-shot task**), or
  - the agent may be teleported to (**continual task**)



## Example: Balancing Cart-pole



**Task:** Make the pole upright for as long as possible.

- Notion of state?  $s = x$  or  $s = (x, \theta)$ ? Neither will yield an MDP definition.
- State:  $s = (x, \dot{x}, \theta, \dot{\theta})$
- Action: Force  $F$
- Reward: If  $|\theta| > \text{th}$ , then 0, else 1. (E.g.,  $\text{th} = 10 \text{ deg}$ )

Once  $|\theta| > \text{th}$ , an episode is terminated, and the pole is put at  $\theta = 0$ .



## Example: Query Optimization

Query Optimization via RL (Marcus et al. (2019); Chen et al. (2023))


**Task:** Find a query plan minimizing long-term execution time

- State: Partial query plan
- Action: Join order decisions
- Reward: negative of execution cost



## Example: RiverSwim

The  $L$ -state RiverSwim MDP



Exercise: Determine

- **State Space:**
- **Action Space:**
- **Reward Function:**
- **Transition Function:**



# Policy



# Policy

When interacting with an MDP, actions are taken according to some **policy**:

Classification of policies:

- deterministic vs. randomized (stochastic)
- stationary vs. history-dependent

	deterministic	randomized
stationary		
history-dependent		



## Stationary Policies

A **stationary deterministic** policy  $\pi$  is a mapping  $\pi : \mathcal{S} \rightarrow \mathcal{A}$ .

- Notation:  $a = \pi(s)$
- $\pi$  prescribes an action with certainty at any state  $s$ , without dependence on past states or actions.

A **stationary randomized** policy  $\pi$  is a mapping  $\pi : \mathcal{S} \rightarrow \Delta(\mathcal{A})$ , where  $\Delta(\mathcal{A})$  denotes the set of probability distributions over  $\mathcal{A}$ .

- Notation:  $a \sim \pi(\cdot|s)$  or  $\pi(a|s)$  denotes the probability of selecting  $a$  in  $s$ :

$$\sum_{a \in \mathcal{A}} \pi(a|s) = 1$$

- At any state  $s$ ,  $\pi$  prescribes a probability distribution over  $\mathcal{A}$ , but without dependence on past states or actions.



## History-dependent Policies

Let  $\mathcal{H}$  the set of all possible histories (trajectories).

A **history-dependent deterministic** policy  $\pi$  is a mapping  $\pi : \mathcal{H} \rightarrow \mathcal{A}$ .

- Notation:  $a = \pi(h_t)$  at time  $t$
- $\pi$  prescribes an action with certainty at any state  $s$ , but depends on past states or actions.

A **history-dependent randomized** policy  $\pi$  is a mapping  $\pi : \mathcal{H} \rightarrow \Delta(\mathcal{A})$ .

- Notation:  $a \sim \pi(\cdot|h_t)$  or  $\pi(a|h_t)$  denotes the probability of selecting  $a$  given history  $h_t$ :

$$\sum_{a \in \mathcal{A}} \pi(a|h_t) = 1, \quad \forall t.$$

- Given any history  $h_t$ ,  $\pi$  prescribes a probability distribution over  $\mathcal{A}$ , arbitrarily depending on past states or actions.



# Policy

	deterministic	randomized
stationary	$\pi : \mathcal{S} \rightarrow \mathcal{A}$	$\pi : \mathcal{S} \rightarrow \Delta(\mathcal{A})$
history-dependent	$\pi : \mathcal{H} \rightarrow \mathcal{A}$	$\pi : \mathcal{H} \rightarrow \Delta(\mathcal{A})$

- $\Pi^{\text{SD}}$ : The set of stationary deterministic policies
- $\Pi^{\text{SR}}$ : The set of stationary randomized policies
- $\Pi^{\text{HD}}$ : The set of history-dependent deterministic policies
- $\Pi^{\text{HR}}$ : The set of history-dependent randomized policies

(i)  $\Pi^{\text{SD}} \subset \Pi^{\text{SR}} \subset \Pi^{\text{HR}}$

(ii)  $\Pi^{\text{SD}} \subset \Pi^{\text{HD}} \subset \Pi^{\text{HR}}$

Notation:


- For  $\pi \in \Pi^{\text{SR}}$ , we write  $a \sim \pi(\cdot|s)$ . Also, given  $f : \mathcal{A} \rightarrow \mathbb{R}$ ,

$$\mathbb{E}_{a \sim \pi(s)}[f(a)] = \sum_{a \in \mathcal{A}} f(a)\pi(a|s)$$



## Policy: Examples

### The $L$ -state RiverSwim MDP



Examples:

- $\pi_1$  : always go right. ( $\pi_1 \in \Pi^{\text{SD}}$ )
- $\pi_2$  : go right w.p. 0.7 and left w.p. 0.3. ( $\pi_2 \in \Pi^{\text{SR}}$ )
- $\pi_3$  : go right if  $s_t \neq s_{t-1}$ , otherwise go left . ( $\pi_3 \in \Pi^{\text{HD}}$ )



## Induced Markov Chains

- Every  $\pi \in \Pi^{\text{SR}}$  induces a **Markov chain** on  $M$ , with transition probability matrix  $P^\pi$  given by:

$$P_{s,s'}^\pi = \sum_{a \in \mathcal{A}} P(s'|s, a)\pi(a|s), \quad s, s' \in \mathcal{S}.$$

- Every  $\pi \in \Pi^{\text{SR}}$  induces a reward vector  $r^\pi \in \mathbb{R}^S$  on  $M$  defined by:

$$r^\pi(s) = \sum_{a \in \mathcal{A}} R(s, a)\pi(a|s), \quad s \in \mathcal{S}.$$

- If  $\pi \in \Pi^{\text{SD}}$ , then  $P_{s,s'}^\pi = P(s'|s, \pi(s))$  and  $r^\pi(s) = R(s, \pi(s))$ .

Every policy  $\pi \in \Pi^{\text{SR}}$  induces a **Markov Reward Process (MRP)** on  $M$ , specified by  $r^\pi$  and  $P^\pi$ .



# Beyond Full Observability




## Partial Observability

- MDPs (and many other decision processes) rest on **full observability** of the state.
- In some applications, the state is **unobservable**, but it can be inferred or estimated via some proxy.
- Some related decision processes: Partially Observable MDP (POMDP), Predictive State Representation (PSR), Regular Decision Process.
- RL under partial observability is much more challenging than in MDPs.

RL under partial observability is beyond the scope of OReL.



## Example



An image from the Malmo platform built for AI experimentation (Photo from (Dann et al., 2018))

- The task is governed by small but **hidden** state-space.
- The agent may infer the current state  $s_t$  via rich sensory observations encoded as  $x_t$ .
- The problem is Markovian w.r.t.  $s$ , not  $x$ .

